

[Peer Assessments \(https://class.coursera.org/rprog-003/human_grading/\)](https://class.coursera.org/rprog-003/human_grading/)

/ Programming Assignment 2: Lexical Scoping

[Help \(https://class.coursera.org/rprog-003/help/peergrading?url=https%3A%2F%2Fclass.coursera.org%2Frprog-003%2Fhuman_grading%2Fview%2Fcourses%2F972138%2Fassessments%2F3%2Fsubmissions\)](https://class.coursera.org/rprog-003/help/peergrading?url=https%3A%2F%2Fclass.coursera.org%2Frprog-003%2Fhuman_grading%2Fview%2Fcourses%2F972138%2Fassessments%2F3%2Fsubmissions)

Submission Phase

1. Do assignment ☐ (/rprog-003/human_grading/view/courses/972138/assessments/3/submissions)
due in 4day 5h

Evaluation Phase

2. Evaluate peers  (/rprog-003/human_grading/view/courses/972138/assessments/3/peerGradingSets)

Results Phase

3. See results  (/rprog-003/human_grading/view/courses/972138/assessments/3/results/mine)

No work was submitted before the submission deadline. You will not be able to evaluate the work of your peers or receive an evaluation.

Introduction

This second programming assignment will require you to write an R function is able to cache potentially time-consuming computations. For example, taking the mean of a numeric vector is typically a fast operation. However, for a very long vector, it may take too long to compute the mean, especially if it has to be computed repeatedly (e.g. in a loop). If the contents of a vector are not changing, it may make sense to cache the value of the mean so that when we need it again, it can be looked up in the cache rather than recomputed. In this Programming Assignment will take advantage of the scoping rules of the R language and how they can be manipulated to preserve state inside of an R object.

Example: Caching the Mean of a Vector

In this example we introduce the `<<-` operator which can be used to assign a value to an object in an environment that is different from the current environment. Below are two functions that are used to create a special object that stores a numeric vector and cache's its mean.

The first function, `makeVector` creates a special "vector", which is really a list containing a function to

1. set the value of the vector

2. get the value of the vector
3. set the value of the mean
4. get the value of the mean

```
makeVector <- function(x = numeric()) {  
  m <- NULL  
  set <- function(y) {  
    x <- y  
    m <- NULL  
  }  
  get <- function() x  
  setmean <- function(mean) m <- mean  
  getmean <- function() m  
  list(set = set, get = get,  
        setmean = setmean,  
        getmean = getmean)  
}
```

The following function calculates the mean of the special "vector" created with the above function. However, it first checks to see if the mean has already been calculated. If so, it gets the mean from the cache and skips the computation. Otherwise, it calculates the mean of the data and sets the value of the mean in the cache via the `setmean` function.

```
cachemean <- function(x, ...) {  
  m <- x$getmean()  
  if(!is.null(m)) {  
    message("getting cached data")  
    return(m)  
  }  
  data <- x$get()  
  m <- mean(data, ...)  
  x$setmean(m)  
  m  
}
```

Assignment: Caching the Inverse of a Matrix

Matrix inversion is usually a costly computation and there may be some benefit to caching the inverse of a matrix rather than compute it repeatedly (there are also alternatives to matrix inversion that we will not discuss here). Your assignment is to write a pair of functions that cache the inverse of a matrix.

Write the following functions:

1. `makeCacheMatrix`: This function creates a special "matrix" object that can cache its inverse.
2. `cacheSolve`: This function computes the inverse of the special "matrix" returned by `makeCacheMatrix` above. If the inverse has already been calculated (and the matrix has not changed), then the `cacheSolve` should retrieve the inverse from the cache.

Computing the inverse of a square matrix can be done with the `solve` function in R. For example, if `x` is a square invertible matrix, then `solve(X)` returns its inverse.

For this assignment, assume that the matrix supplied is always invertible.

In order to complete this assignment, you must do the following:

1. Fork the GitHub repository containing the stub R files at <https://github.com/rdpeng/ProgrammingAssignment2> (<https://github.com/rdpeng/ProgrammingAssignment2>) to create a copy under your own account.
2. Clone your forked GitHub repository to your computer so that you can edit the files locally on your own machine.
3. Edit the R file contained in the git repository and place your solution in that file (please do not rename the file).
4. Commit your completed R file into YOUR git repository and push your git branch to the GitHub repository under your account.
5. Submit to Coursera the URL to your GitHub repository that contains the completed R code for the assignment.

Grading

This assignment will be graded via peer assessment.

Please submit the URL pointing to your GitHub repository containing the completed R code for this assignment.

Overall evaluation/feedback

Note: this section can only be filled out during the evaluation phase.

Was a valid GitHub URL containing a git repository submitted?

0 points: A valid GitHub URL was NOT submitted (or URL is broken)
1 point: The submitted URL points to a GitHub repository

Does the GitHub repository contain at least one commit beyond the original fork?

0 points: No, there are no commits beyond the original fork

1 point: Yes, there is at least one commit beyond the original fork

Does the GitHub repository contain an R file containing code implementing the completed assignment?

NOTE: There is no need to run the code here, but rather you should visually inspect the R file in the GitHub repository and check to see that there is code there beyond the original stub that was committed there by the instructor. Do not attempt to judge whether the code is correct or not.

0 points: No, the R file does not contain any code implementing the completed assignment

1 point: The R file contains code implementing a partially completed assignment

2 points: The R file contains code implementing the complete assignment

Does the R file containing the code have any comments explaining what the code does?

NOTE: The makeCacheMatrix and cacheSolve functions should both be documented with explanatory comments. There maybe other functions in the R file but they do not need to be commented.

0 points: There are no explanatory comments in the R file

1 points: One of the functions has some corresponding explanatory comments in the R file

2 points: Both functions have corresponding explanatory comments in the R file

Does the R code implementing the 'makeCacheMatrix' function appear to be correct, to the best of your ability to judge?

NOTE: Do not run the R code on your own computer. Please examine the R code and

determine to the best of your ability, whether the solution presented appears to match the requirements of the assignment. A correct solution (as best you can determine) gets 1 point and a solution that is well-written and easy to read gets an additional 1 point.

0 points: The solution contains at least one identifiable problem

1 point: The solution appears to be correct

2 points: The solution appears to be correct and is written in a well-formatted and easily readable style

Does the R code implementing the 'cacheSolve' function appear to be correct, to the best of your ability to judge?

NOTE: Do not run the R code on your own computer. Please examine the R code and determine to the best of your ability, whether the solution presented appears to match the requirements of the assignment. A correct solution (as best you can determine) gets 1 point and a solution that is well-written and easy to read gets an additional 1 point.

0 points: The solution contains at least one identifiable problem

1 point: The solution appears to be correct

2 points: The solution appears to be correct and is written in a well-formatted and easily readable style